

# Assignment\_GenAI: Detailed Report

## Project Overview

This project is an AI-powered **PDF-based Question Answering system**. It allows users to:

1. Upload a PDF.
  2. Ask questions related to that PDF.
  3. Get answers generated using **Google's Large Language Model (LLM)**, with relevant context retrieved via **SentenceTransformer** and **ChromaDB** (vector store).
- 

## Components Involved

### 1. SentenceTransformer

- Used to convert chunks of the PDF content into embeddings.
- Enables semantic search for relevant content to the question.

### 2. Google LLM (**get\_answer**)

- Uses the Google Generative Language API to generate answers based on context returned by the embedding search.
  - Context-aware QA.
- 

## Endpoints in **main.py**

### **POST /upload\_pdf/**

- **Purpose:** Uploads and processes a PDF.
- **Inputs:** PDF file and **pdf\_name**.
- **Process:**

- Saves the file.
- Calls `process_pdf()` to extract content and store it in ChromaDB (vector DB).
- **Output:** JSON success message.

## POST `/ask_question/`

- **Purpose:** Answers a question related to a specific PDF.
  - **Inputs:**
    - `question`: the user's query.
    - `pdf_name`: the name of the already-uploaded PDF.
  - **Process:**
    - Calls `search_chroma()` to find relevant chunks.
    - Calls `get_answer()` to generate an answer using Google's LLM.
  - **Output:** The answer string.
- 

Next, I'll analyze the logic in `pdf_processor.py` and `qa_engine.py` to include descriptions of all major functions and files in the final report.

Here's the breakdown of the key files and functions based on your `utils` folder:

---

### `pdf_processor.py`

**Purpose:** Extract text from PDFs, chunk the content, convert it into embeddings, and store it in ChromaDB.

#### **Key Components:**

- **PdfReader:** Reads the uploaded PDF file.
- **RecursiveCharacterTextSplitter:** Splits the text into manageable chunks.

- **SentenceTransformerEmbeddings:**
  - Model used: `all-MiniLM-L6-v2`
  - Converts chunks into dense vector embeddings.

#### ⚙️ Key Function:

##### `process_pdf(pdf_path, pdf_name)`

- Reads and extracts text from the PDF.
- Chunks the text.
- Creates vector embeddings using SentenceTransformer.
- Stores them in ChromaDB under the collection named after the PDF.

#### 🧠 `qa_engine.py`

**Purpose:** Performs semantic search and answers user questions using Google Gemini.

#### 🔧 Key Components:

- **Google Generative AI (`gemini-1.5-pro-002`)** via LangChain.
- **Chroma:** Retrieves similar content using embeddings.
- **SentenceTransformer:** Reused from `pdf_processor`.

#### 🔍 Functions:

##### `search_chroma(question, pdf_name)`

- Uses the vector store to find the top relevant chunks from the specified PDF based on the question.
- Returns those chunks as context.

##### `get_answer(context, question)`

- Sends both the `context` and `question` to Google Gemini model.
- Returns a high-quality natural language answer.